

The Working Pack

Everything the team works from lives in one folder: gauntlet - day/. The Submissions Book says what's claimed; the six-sections handout says what a validation report is made of; this page says what's in the folder and why. The build target: one pipeline, run unchanged over all seven submissions, returning PASS, SUSPECT with conditions attached, or FAIL — with the evidence behind every line.

data/ — the market as it actually traded

Six tables, all timestamps UTC, schema in DATA.md: **book_snapshots** (the intraday order book, three price/size levels per side, every 15 minutes for the last 10 hours before gate), **forecasts** (every wind/solar/load vintage with its issue timestamp — what was knowable at any moment is exactly the rows issued by then), **products**, **da_auction**, **actuals**, and **outages** (announced planned unavailability, published day-ahead). There is no mid column: mid is a derived convenience, $(bid + ask) / 2$ — not a price anyone quoted. The claimed numbers only reproduce at mid; treat that as a finding, not a convention.

blotters/ — the behavioural record

One CSV per submission, every trade: `exec_ts`, `product_delivery`, `side`, `qty_mw`, `price`. One schema for all seven, so anything built runs on all of them unchanged. Every position is opened and closed through trades, so a product's P&L is sells minus buys at whatever price column values the trades. The blotter is what makes records-based validation possible without ever seeing code.

registry/ — the declarations

One YAML per submission: params as declared, declared trial count, dev window, data dependencies, and the claimed headline stats. **None of it is verified** — the registry is the submitter speaking, and checking the statements against the records is the pipeline's job. Luck computes on the declared numbers (deflated Sharpe reads `declared_trials` from here); Lineage checks whether the declarations are true. Each file ends with an empty `validation:` block — verdict, conditions, warranty, evidence — that the pipeline fills. `loader.py` loads all seven.

starter/ — the scaffold, nothing more

`00_start_here.ipynb` loads everything and closes the consistency loop: every blotter, repriced at the book mid over its declared window, reproduces its claimed P&L and Sharpe to rounding — for all seven, including the broken ones. That's where validation starts, not where it ends. `repricer.py` does mid repricing only, deliberately: repricing at prices you'd actually get — crossing the spread, walking the book — is the team's to build, and most of the Friction section stands on it. `report.py` is the contract:

```
run(registry_entry, blotter, market) -> Report
Report.sections = [luck, lineage, friction, shelf_life, evidence, warranty]
Report.verdict: any section FAIL => FAIL; any SUSPECT => SUSPECT + conditions
```

No validation logic is pre-written. The six sections say what to compute; how is the day.

docs/ — the two handouts

The Submissions Book and the six-sections handout, so the folder is self-contained.

What is deliberately not in the folder

No strategy source code and no callables — a validator gets records, and records are what these verdicts must stand on. No generator, no answer key. The truth about the seven, and six further months of market data, stay sealed until the reveal: verdicts are judged against the truth, warranties against the future.